

Questions (69)

- [175. Combine Two Tables](#)
- [176. Second Highest Salary](#)
- [177. Nth Highest Salary](#)
- [178. Rank Scores](#)
- [180. Consecutive Numbers](#)
- [181. Employees Earning More Than Their Managers](#)
- [182. Duplicate Emails](#)
- [183. Customers Who Never Order](#)
- [196. Delete Duplicate Emails](#)
- [197. Rising Temperature](#)
- [511. Game Play Analysis I](#)
- [512. Game Play Analysis II](#)
- [534. Game Play Analysis III](#)
- [570. Managers with at Least 5 Direct Reports](#)
- [577. Employee Bonus](#)
- [584. Find Customer Referee](#)
- [586. Customer Placing the Largest Number of Orders](#)
- [595. Big Countries](#)
- [596. Classes More Than 5 Students](#)
- [597. Friend Requests I: Overall Acceptance Rate](#)
- [603. Consecutive Available Seats](#)
- [607. Sales Person](#)
- [608. Tree Node](#)
- [610. Triangle Judgement](#)
- [612. Shortest Distance in a Plane](#)
- [613. Shortest Distance in a Line](#)

[619. Biggest Single Number](#)

[620. Not Boring Movies](#)

[626. Exchange Seats](#)

[627. Swap Salary](#)

[1045. Customers Who Bought All Products](#)

[1050. Actors and Directors Who Cooperated At Least Three Times](#)

[1068. Product Sales Analysis I](#)

[1069. Product Sales Analysis II](#)

[1070. Product Sales Analysis III](#)

[1075. Project Employees I](#)

[1076. Project Employees II](#)

[1077. Project Employees III](#)

[1082. Sales Analysis I](#)

[1083. Sales Analysis II](#)

[1084. Sales Analysis III](#)

[1112. Highest Grade For Each Student](#)

[1113. Reported Posts](#)

[1126. Active Businesses](#)

[1141. User Activity for the Past 30 Days I](#)

[1142. User Activity for the Past 30 Days II](#)

[1148. Article Views I](#)

[1164. Product Price at a Given Date](#)

[1173. Immediate Food Delivery I](#)

[1174. Immediate Food Delivery II](#)

[1179. Reformat Department Table](#)

[1193. Monthly Transactions I](#)

[1204. Last Person to Fit in the Elevator](#)

[1211. Queries Quality and Percentage](#)

[1212. Team Scores in Football Tournament](#)
[1225. Report Contiguous Dates](#)
[1241. Number of Comments per Post](#)
[1251. Average Selling Price](#)
[1264. Page Recommendations](#)
[1270. All People Report to the Given Manager](#)
[1280. Students and Examinations](#)
[1285. Find the Start and End Number of Continuous Ranges](#)
[1294. Weather Type in Each Country](#)
[1303. Find the Team Size](#)
[1308. Running Total for Different Genders](#)
[1321. Restaurant Growth](#)
[1322. Ads Performance](#)
[1327. List the Products Ordered in a Period](#)
[1336. Number of Transactions per Visit](#)

Easy (43)

[175. Combine Two Tables](#)
[181. Employees Earning More Than Their Managers](#)
[182. Duplicate Emails](#)
[183. Customers Who Never Order](#)
[196. Delete Duplicate Emails](#)
[197. Rising Temperature](#)
[511. Game Play Analysis I](#)
[512. Game Play Analysis II](#)
[577. Employee Bonus](#)
[584. Find Customer Referee](#)
[586. Customer Placing the Largest Number of Orders](#)

[595. Big Countries](#)

[596. Classes More Than 5 Students](#)

[597. Friend Requests I: Overall Acceptance Rate](#)

[603. Consecutive Available Seats](#)

[607. Sales Person](#)

[610. Triangle Judgement](#)

[613. Shortest Distance in a Line](#)

[619. Biggest Single Number](#)

[620. Not Boring Movies](#)

[627. Swap Salary](#)

[1050. Actors and Directors Who Cooperated At Least Three Times](#)

[1068. Product Sales Analysis I](#)

[1069. Product Sales Analysis II](#)

[1075. Project Employees I](#)

[1076. Project Employees II](#)

[1082. Sales Analysis I](#)

[1083. Sales Analysis II](#)

[1084. Sales Analysis III](#)

[1113. Reported Posts](#)

[1141. User Activity for the Past 30 Days I](#)

[1142. User Activity for the Past 30 Days II](#)

[1148. Article Views I](#)

[1173. Immediate Food Delivery I](#)

[1179. Reformat Department Table](#)

[1211. Queries Quality and Percentage](#)

[1241. Number of Comments per Post](#)

[1251. Average Selling Price](#)

[1280. Students and Examinations](#)

[1294. Weather Type in Each Country](#)

[1303. Find the Team Size](#)

[1322. Ads Performance](#)

[1327. List the Products Ordered in a Period](#)

Medium (24)

[176. Second Highest Salary](#)

[177. Nth Highest Salary](#)

[178. Rank Scores](#)

[180. Consecutive Numbers](#)

[534. Game Play Analysis III](#)

[570. Managers with at Least 5 Direct Reports](#)

[608. Tree Node](#)

[612. Shortest Distance in a Plane](#)

[626. Exchange Seats](#)

[1045. Customers Who Bought All Products](#)

[1070. Product Sales Analysis III](#)

[1077. Project Employees III](#)

[1112. Highest Grade For Each Student](#)

[1126. Active Businesses](#)

[1164. Product Price at a Given Date](#)

[1174. Immediate Food Delivery II](#)

[1193. Monthly Transactions I](#)

[1204. Last Person to Fit in the Elevator](#)

[1212. Team Scores in Football Tournament](#)

[1264. Page Recommendations](#)

[1270. All People Report to the Given Manager](#)

[1285. Find the Start and End Number of Continuous Ranges](#)

[1308. Running Total for Different Genders](#)

[1321. Restaurant Growth](#)

Hard (2)

[1225. Report Contiguous Dates](#)

[1336. Number of Transactions per Visit](#)

175. Combine Two Tables

Table: Person

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| PersonId    | int    |
| FirstName    | varchar|
| LastName     | varchar|
+-----+-----+
```

PersonId is the primary key column for this table.

Table: Address

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| AddressId    | int    |
| PersonId     | int    |
| City         | varchar|
| State        | varchar|
+-----+-----+
```

AddressId is the primary key column for this table.

Write a SQL query for a report that provides the following information for each person in the Person table, regardless if there is an address for each of those people:

FirstName, LastName, City, State

176. Second Highest Salary

Write a SQL query to get the second highest salary from the Employee table.

```
+-----+-----+
| Id | Salary |
+-----+-----+
| 1  | 100    |
| 2  | 200    |
| 3  | 300    |
+-----+-----+
```

For example, given the above Employee table, the query should return 200 as the second highest salary. If there is no second highest salary, then the query should return null.

```
+-----+
| SecondHighestSalary |
+-----+
| 200                  |
+-----+
```

177. Nth Highest Salary

Write a SQL query to get the nth highest salary from the Employee table.

```
+-----+-----+
| Id | Salary |
+-----+-----+
| 1  | 100    |
| 2  | 200    |
| 3  | 300    |
+-----+-----+
```

For example, given the above Employee table, the nth highest salary where n = 2 is 200. If there is no nth highest salary, then the query should return null.

```
+-----+
| getNthHighestSalary(2) |
+-----+
| 200                    |
+-----+
```

178. Rank Scores

Write a SQL query to rank scores. If there is a tie between two scores, both should have the same ranking. Note that after a tie, the next ranking number should be the next consecutive integer value. In other words, there should be no "holes" between ranks.

```
+----+-----+
| Id | Score |
+----+-----+
| 1  | 3.50  |
| 2  | 3.65  |
| 3  | 4.00  |
| 4  | 3.85  |
| 5  | 4.00  |
| 6  | 3.65  |
+----+-----+
```

For example, given the above Scores table, your query should generate the following report (order by highest score):

```
+-----+-----+
| Score | Rank |
+-----+-----+
| 4.00  | 1    |
| 4.00  | 1    |
| 3.85  | 2    |
| 3.65  | 3    |
| 3.65  | 3    |
| 3.50  | 4    |
+-----+-----+
```

180. Consecutive Numbers

Write a SQL query to find all numbers that appear at least three times consecutively.

```
+----+-----+
| Id | Num |
+----+-----+
| 1  | 1   |
| 2  | 1   |
| 3  | 1   |
| 4  | 2   |
| 5  | 1   |
| 6  | 2   |
| 7  | 2   |
+----+-----+
```


For example, given the above Logs table, 1 is the only number that appears consecutively for at least three times.

```
+-----+
| ConsecutiveNums |
+-----+
| 1                |
+-----+
```

181. Employees Earning More Than Their Managers

The Employee table holds all employees including their managers. Every employee has an Id, and there is also a column for the manager Id.

```
+----+-----+-----+-----+
| Id | Name  | Salary | ManagerId |
+----+-----+-----+-----+
| 1  | Joe   | 70000  | 3          |
| 2  | Henry | 80000  | 4          |
| 3  | Sam   | 60000  | NULL       |
| 4  | Max   | 90000  | NULL       |
+----+-----+-----+-----+
```

Given the Employee table, write a SQL query that finds out employees who earn more than their managers. For the above table, Joe is the only employee who earns more than his manager.

```
+-----+
| Employee |
+-----+
| Joe      |
+-----+
```

182. Duplicate Emails

Write a SQL query to find all duplicate emails in a table named Person.

```
+----+-----+
| Id | Email    |
+----+-----+
| 1  | a@b.com  |
| 2  | c@d.com  |
| 3  | a@b.com  |
+----+-----+
```

For example, your query should return the following for the above table:

```
+-----+
| Email  |
+-----+
| a@b.com |
+-----+
```

Note: All emails are in lowercase.

183. Customers Who Never Order

Suppose that a website contains two tables, the Customers table and the Orders table. Write a SQL query to find all customers who never order anything.

Table: Customers.

```
+-----+-----+
| Id | Name |
+-----+-----+
| 1  | Joe  |
| 2  | Henry |
| 3  | Sam  |
| 4  | Max  |
+-----+-----+
```

Table: Orders.

```
+-----+-----+
| Id | CustomerId |
+-----+-----+
| 1  | 3          |
| 2  | 1          |
+-----+-----+
```

Using the above tables as example, return the following:

```
+-----+
| Customers |
+-----+
| Henry     |
| Max       |
+-----+
```

196. Delete Duplicate Emails

Write a SQL query to delete all duplicate email entries in a table named Person, keeping only unique emails based on its smallest Id.

```
+-----+-----+
| Id | Email |
+-----+-----+
| 1 | john@example.com |
| 2 | bob@example.com |
| 3 | john@example.com |
+-----+-----+
```

Id is the primary key column for this table.

For example, after running your query, the above Person table should have the following rows:

```
+-----+-----+
| Id | Email |
+-----+-----+
| 1 | john@example.com |
| 2 | bob@example.com |
+-----+-----+
```

Note:

Your output is the whole Person table after executing your sql. Use delete statement.

197. Rising Temperature

Given a Weather table, write a SQL query to find all dates' Ids with higher temperature compared to its previous (yesterday's) dates.

```
+-----+-----+-----+
| Id(INT) | RecordDate (DATE) | Temperature (INT) |
+-----+-----+-----+
| 1 | 2015-01-01 | 10 |
| 2 | 2015-01-02 | 25 |
| 3 | 2015-01-03 | 20 |
| 4 | 2015-01-04 | 30 |
+-----+-----+-----+
```

For example, return the following Ids for the above Weather table:

```
+-----+
| Id |
+-----+
| 2 |
| 4 |
```

+-----+

511. Game Play Analysis I

Table: Activity

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player_id, event_date) is the primary key of this table.

This table shows the activity of players of some game.

Each row is a record of a player who logged in and played a number of games (possibly 0) before logging out on some day using some device.

Write an SQL query that reports the first login date for each player.

The query result format is in the following example:

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-05-02	6
2	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Result table:

player_id	first_login
1	2016-03-01
2	2017-06-25
3	2016-03-02

512. Game Play Analysis II

Table: Activity

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player_id, event_date) is the primary key of this table.

This table shows the activity of players of some game.

Each row is a record of a player who logged in and played a number of games (possibly 0) before logging out on some day using some device.

Write a SQL query that reports the device that is first logged in for each player.

The query result format is in the following example:

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-05-02	6
2	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Result table:

player_id	device_id
1	2
2	3
3	1

534. Game Play Analysis III

Table: Activity

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player_id, event_date) is the primary key of this table.

This table shows the activity of players of some game.

Each row is a record of a player who logged in and played a number of games (possibly 0) before logging out on some day using some device.

Write an SQL query that reports for each player and date, how many games played so far by the player. That is, the total number of games played by the player until that date. Check the example for clarity.

The query result format is in the following example:

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-05-02	6
1	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Result table:

player_id	event_date	games_played_so_far
1	2016-03-01	5
1	2016-05-02	11
1	2017-06-25	12
3	2016-03-02	0
3	2018-07-03	5

For the player with id 1, $5 + 6 = 11$ games played by 2016-05-02, and $5 + 6 + 1 = 12$ games played by 2017-06-25.

For the player with id 3, $0 + 5 = 5$ games played by 2018-07-03.

Note that for each player we only care about the days when the player logged in.

570. Managers with at Least 5 Direct Reports

The Employee table holds all employees including their managers. Every employee has an Id, and there is also a column for the manager Id.

Id	Name	Department	ManagerId
101	John	A	null
102	Dan	A	101
103	James	A	101
104	Amy	A	101
105	Anne	A	101
106	Ron	B	101

Given the Employee table, write a SQL query that finds out managers with at least 5 direct reports. For the above table, your SQL query should return:

Name
John

Note:

No one would report to himself.

577. Employee Bonus

Select all employee's name and bonus whose bonus is < 1000.

Table:Employee

empId	name	supervisor	salary
1	John	3	1000
2	Dan	3	2000
3	Brad	null	4000
4	Thomas	3	4000

```
+-----+-----+-----+-----+
empId is the primary key column for this table.
Table: Bonus
```

```
+-----+-----+
| empId | bonus |
+-----+-----+
| 2     | 500   |
| 4     | 2000  |
+-----+-----+
```

```
empId is the primary key column for this table.
Example output:
```

```
+-----+-----+
| name  | bonus |
+-----+-----+
| John  | null  |
| Dan   | 500   |
| Brad  | null  |
+-----+-----+
```

584. Find Customer Referee

Given a table customer holding customers information and the referee.

```
+-----+-----+-----+
| id  | name | referee_id|
+-----+-----+-----+
| 1   | Will |          NULL |
| 2   | Jane |          NULL |
| 3   | Alex |           2 |
| 4   | Bill |          NULL |
| 5   | Zack |           1 |
| 6   | Mark |           2 |
+-----+-----+-----+
```

Write a query to return the list of customers NOT referred by the person with id '2'.

For the sample data above, the result is:

```
+-----+
| name |
+-----+
| Will |
| Jane |
```



```
| Bill |
| Zack |
+-----+
```

586. Customer Placing the Largest Number of Orders

Query the customer_number from the orders table for the customer who has placed the largest number of orders.

It is guaranteed that exactly one customer will have placed more orders than any other customer.

The orders table is defined as follows:

Column	Type
order_number (PK)	int
customer_number	int
order_date	date
required_date	date
shipped_date	date
status	char(15)
comment	char(200)

Sample Input

order_number	customer_number	order_date	required_date	shipped_date	status	comment
1	1	2017-04-09	2017-04-13	2017-04-12	Closed	
2	2	2017-04-15	2017-04-20	2017-04-18	Closed	
3	3	2017-04-16	2017-04-25	2017-04-20	Closed	
4	3	2017-04-18	2017-04-28	2017-04-25	Closed	

Sample Output

```
| customer_number |
|-----|
| 3              |
```

Explanation

The customer with number '3' has two orders, which is greater than either customer '1' or '2' because each of them only has one order.
 So the result is customer_number '3'.
 Follow up: What if more than one customer have the largest number of orders, can you find all the customer_number in this case?

595. Big Countries

There is a table World

name	continent	area	population	gdp
Afghanistan	Asia	652230	25500100	20343000
Albania	Europe	28748	2831741	12960000
Algeria	Africa	2381741	37100000	188681000
Andorra	Europe	468	78115	3712000
Angola	Africa	1246700	20609294	100990000

A country is big if it has an area of bigger than 3 million square km or a population of more than 25 million.

Write a SQL solution to output big countries' name, population and area.

For example, according to the above table, we should output:

name	population	area
Afghanistan	25500100	652230
Algeria	37100000	2381741

596. Classes More Than 5 Students

There is a table courses with columns: student and class

Please list out all classes which have more than or equal to 5 students.

For example, the table:

student	class
A	Math
B	English
C	Math
D	Biology
E	Math
F	Computer
G	Math
H	Math
I	Math

Should output:

class
Math

Note:

The students should not be counted duplicate in each course.

597. Friend Requests I: Overall Acceptance Rate

In social network like Facebook or Twitter, people send friend requests and accept others' requests as well. Now given two tables as below:

Table: friend_request

sender_id	send_to_id	request_date
1	2	2016_06-01
1	3	2016_06-01
1	4	2016_06-01
2	3	2016_06-02
3	4	2016-06-09

Table: request_accepted

requester_id	accepter_id	accept_date
1	2	2016_06-03
1	3	2016-06-08
2	3	2016-06-08
3	4	2016-06-09
3	4	2016-06-10

Write a query to find the overall acceptance rate of requests rounded to 2 decimals, which is the number of acceptance divide the number of requests.

For the sample data above, your query should return the following result.

accept_rate
0.80

Note:

The accepted requests are not necessarily from the table friend_request. In this case, you just need to simply count the total accepted requests (no matter whether they are in the original requests), and divide it by the number of requests to get the acceptance rate.

It is possible that a sender sends multiple requests to the same receiver, and a request could be accepted more than once. In this case, the 'duplicated' requests or acceptances are only counted once.

If there is no requests at all, you should return 0.00 as the accept_rate.

Explanation: There are 4 unique accepted requests, and there are 5 requests in total. So the rate is 0.80.

Follow-up:

Can you write a query to return the accept rate but for every month?

How about the cumulative accept rate for every day?

603. Consecutive Available Seats

Several friends at a cinema ticket office would like to reserve consecutive available seats.

Can you help to query all the consecutive available seats order by the seat_id using the following cinema table?

seat_id	free
1	1
2	0
3	1
4	1
5	1

Your query should return the following result for the sample case above.

seat_id
3
4
5

Note:

The seat_id is an auto increment int, and free is bool ('1' means free, and '0' means occupied.).

Consecutive available seats are more than 2(inclusive) seats consecutively available.

607. Sales Person

Given three tables: salesperson, company, orders.

Output all the names in the table salesperson, who didn't have sales to company 'RED'.

Example

Input

Table: salesperson

sales_id	name	salary	commission_rate	hire_date
1	John	100000	6	4/1/2006
2	Amy	120000	5	5/1/2010
3	Mark	65000	12	12/25/2008
4	Pam	25000	25	1/1/2005

5	Alex	50000	10	2/3/2007
---	------	-------	----	----------

The table salesperson holds the salesperson information. Every salesperson has a sales_id and a name.

Table: company

com_id	name	city
1	RED	Boston
2	ORANGE	New York
3	YELLOW	Boston
4	GREEN	Austin

The table company holds the company information. Every company has a com_id and a name.

Table: orders

order_id	order_date	com_id	sales_id	amount
1	1/1/2014	3	4	100000
2	2/1/2014	4	5	5000
3	3/1/2014	1	1	50000
4	4/1/2014	1	4	25000

The table orders holds the sales record information, salesperson and customer company are represented by sales_id and com_id.

output

name
Amy
Mark
Alex

Explanation

According to order '3' and '4' in table orders, it is easy to tell only salesperson 'John' and 'Alex' have sales to company 'RED', so we need to output all the other names in table salesperson.

608. Tree Node

Given a table tree, id is identifier of the tree node and p_id is its parent node's id.

```
+-----+-----+
| id | p_id |
+-----+-----+
| 1  | null  |
| 2  | 1     |
| 3  | 1     |
| 4  | 2     |
| 5  | 2     |
+-----+-----+
```

Each node in the tree can be one of three types:

Leaf: if the node is a leaf node.

Root: if the node is the root of the tree.

Inner: If the node is neither a leaf node nor a root node.

Write a query to print the node id and the type of the node. Sort your output by the node id. The result for the above sample is:

```
+-----+-----+
| id | Type |
+-----+-----+
| 1  | Root |
| 2  | Inner|
| 3  | Leaf |
| 4  | Leaf |
| 5  | Leaf |
+-----+-----+
```


Explanation

Node '1' is root node, because its parent node is NULL and it has child node '2' and '3'.

Node '2' is inner node, because it has parent node '1' and child node '4' and '5'.

Node '3', '4' and '5' is Leaf node, because they have parent node and they don't have child node.

And here is the image of the sample tree as below:



Note

If there is only one node on the tree, you only need to output its root attributes.

610. Triangle Judgement

A pupil Tim gets homework to identify whether three line segments could possibly form a triangle.

However, this assignment is very heavy because there are hundreds of records to calculate.

Could you help Tim by writing a query to judge whether these three sides can form a triangle, assuming table triangle holds the length of the three sides x, y and z.

```

| x | y | z |
|---|---|---|
| 13 | 15 | 30 |
| 10 | 20 | 15 |
  
```

For the sample data above, your query should return the follow result:

```

| x | y | z | triangle |
|---|---|---|-----|
| 13 | 15 | 30 | No        |
| 10 | 20 | 15 | Yes       |
  
```

612. Shortest Distance in a Plane

Table point_2d holds the coordinates (x,y) of some unique points (more than two) in a plane.

Write a query to find the shortest distance between these points rounded to 2 decimals.

x	y
-1	-1
0	0
-1	-2

The shortest distance is 1.00 from point (-1,-1) to (-1,2). So the output should be:

shortest
1.00

Note: The longest distance among all the points are less than 10000.

613. Shortest Distance in a Line

Table point holds the x coordinate of some points on x-axis in a plane, which are all integers.

Write a query to find the shortest distance between two points in these points.

x
-1
0
2

The shortest distance is '1' obviously, which is from point '-1' to '0'. So the output is as below:

shortest
1

Note: Every point is unique, which means there is no duplicates in table point.

Follow-up: What if all these points have an id and are arranged from the left most to the right most of x axis?

619. Biggest Single Number

Table my_numbers contains many numbers in column num including duplicated ones.

Can you write a SQL query to find the biggest number, which only appears once.

```
+----+
|num|
+----+
| 8 |
| 8 |
| 3 |
| 3 |
| 1 |
| 4 |
| 5 |
| 6 |
```

For the sample data above, your query should return the following result:

```
+----+
|num|
+----+
| 6 |
```

Note:

If there is no such number, just output null.

620. Not Boring Movies

X city opened a new cinema, many people would like to go to this cinema. The cinema also gives out a poster indicating the movies' ratings and descriptions.

Please write a SQL query to output movies with an odd numbered ID and a description that is not 'boring'. Order the result by rating.

For example, table cinema:

id	movie	description	rating
1	War	great 3D	8.9
2	Science	fiction	8.5
3	irish	boring	6.2
4	Ice song	Fantasy	8.6
5	House card	Interesting	9.1

For the example above, the output should be:

id	movie	description	rating
5	House card	Interesting	9.1
1	War	great 3D	8.9

626. Exchange Seats

Mary is a teacher in a middle school and she has a table seat storing students' names and their corresponding seat ids.

The column id is continuous increment.

Mary wants to change seats for the adjacent students.

Can you write a SQL query to output the result for Mary?

id	student
1	Abbot
2	Doris
3	Emerson
4	Green
5	Jeames

For the sample input, the output is:

id	student
1	Doris
2	Abbot
3	Green
4	Emerson
5	Jeames

Note:

If the number of students is odd, there is no need to change the last one's seat.

627. Swap Salary

Given a table salary, such as the one below, that has m=male and f=female values. Swap all f and m values (i.e., change all f values to m and vice versa) with a single update statement and no intermediate temp table.

Note that you must write a single update statement, DO NOT write any select statement for this problem.

Example:

id	name	sex	salary
1	A	m	2500
2	B	f	1500
3	C	m	5500
4	D	f	500

After running your update statement, the above salary table should have the following rows:

id	name	sex	salary
1	A	f	2500
2	B	m	1500
3	C	f	5500
4	D	m	500

1045. Customers Who Bought All Products

Table: Customer

+-----+-----+		
Column Name	Type	
+-----+-----+		
customer_id	int	
product_key	int	
+-----+-----+		

product_key is a foreign key to Product table.

Table: Product

+-----+-----+		
Column Name	Type	
+-----+-----+		
product_key	int	
+-----+-----+		

product_key is the primary key column for this table.

Write an SQL query for a report that provides the customer ids from the Customer table that bought all the products in the Product table.

For example:

Customer table:

+-----+-----+		
customer_id	product_key	
+-----+-----+		
1	5	
2	6	
3	5	
3	6	
1	6	
+-----+-----+		

Product table:

+-----+		
product_key		
+-----+		
5		
6		
+-----+		

Result table:

+-----+		
customer_id		
+-----+		

```
| 1          |
| 3          |
+-----+
```

The customers who bought all the products (5 and 6) are customers with id 1 and 3.

1050. Actors and Directors Who Cooperated At Least Three Times

Table: ActorDirector

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| actor_id    | int    |
| director_id | int    |
| timestamp   | int    |
+-----+-----+
```

timestamp is the primary key column for this table.

Write a SQL query for a report that provides the pairs (actor_id, director_id) where the actor have cooperated with the director at least 3 times.

Example:

ActorDirector table:

```
+-----+-----+-----+
| actor_id | director_id | timestamp |
+-----+-----+-----+
| 1        | 1           | 0         |
| 1        | 1           | 1         |
| 1        | 1           | 2         |
| 1        | 2           | 3         |
| 1        | 2           | 4         |
| 2        | 1           | 5         |
| 2        | 1           | 6         |
+-----+-----+-----+
```

Result table:

```
+-----+-----+
| actor_id | director_id |
+-----+-----+
```

```
| 1          | 1          |
+-----+-----+
```

The only pair is (1, 1) where they cooperated exactly 3 times.

1068. Product Sales Analysis I

Table: Sales

```
+-----+-----+
| Column Name | Type  |
+-----+-----+
| sale_id     | int   |
| product_id  | int   |
| year        | int   |
| quantity    | int   |
| price       | int   |
+-----+-----+
```

(sale_id, year) is the primary key of this table.

product_id is a foreign key to Product table.

Note that the price is per unit.

Table: Product

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| product_id  | int    |
| product_name | varchar|
+-----+-----+
```

product_id is the primary key of this table.

Write an SQL query that reports all product names of the products in the Sales table along with their selling year and price.

For example:

Sales table:

```
+-----+-----+-----+-----+-----+
| sale_id | product_id | year | quantity | price |
+-----+-----+-----+-----+-----+
| 1       | 100        | 2008 | 10        | 5000  |
| 2       | 100        | 2009 | 12        | 5000  |
| 7       | 200        | 2011 | 15        | 9000  |
```

```
+-----+-----+-----+-----+-----+
```

Product table:

```
+-----+-----+
| product_id | product_name |
+-----+-----+
| 100        | Nokia       |
| 200        | Apple        |
| 300        | Samsung      |
+-----+-----+
```

Result table:

```
+-----+-----+-----+
| product_name | year  | price |
+-----+-----+-----+
| Nokia       | 2008  | 5000  |
| Nokia       | 2009  | 5000  |
| Apple        | 2011  | 9000  |
+-----+-----+-----+
```

1069. Product Sales Analysis II

Table: Sales

```
+-----+-----+
| Column Name | Type  |
+-----+-----+
| sale_id     | int   |
| product_id  | int   |
| year        | int   |
| quantity    | int   |
| price       | int   |
+-----+-----+
```

sale_id is the primary key of this table.

product_id is a foreign key to Product table.

Note that the price is per unit.

Table: Product

```
+-----+-----+
| Column Name | Type  |
+-----+-----+
| product_id  | int   |
| product_name | varchar |
+-----+-----+
```

product_id is the primary key of this table.

Write an SQL query that reports the total quantity sold for every product id.

The query result format is in the following example:

Sales table:

sale_id	product_id	year	quantity	price
1	100	2008	10	5000
2	100	2009	12	5000
7	200	2011	15	9000

Product table:

product_id	product_name
100	Nokia
200	Apple
300	Samsung

Result table:

product_id	total_quantity
100	22
200	15

1070. Product Sales Analysis III

Table: Sales

Column Name	Type
sale_id	int
product_id	int
year	int
quantity	int
price	int

sale_id is the primary key of this table.
product_id is a foreign key to Product table.
Note that the price is per unit.
Table: Product

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| product_id  | int    |
| product_name | varchar |
+-----+-----+
```

product_id is the primary key of this table.

Write an SQL query that selects the product id, year, quantity, and price for the first year of every product sold.

The query result format is in the following example:

Sales table:

```
+-----+-----+-----+-----+-----+
| sale_id | product_id | year | quantity | price |
+-----+-----+-----+-----+-----+
| 1       | 100        | 2008 | 10       | 5000  |
| 2       | 100        | 2009 | 12       | 5000  |
| 7       | 200        | 2011 | 15       | 9000  |
+-----+-----+-----+-----+-----+
```

Product table:

```
+-----+-----+
| product_id | product_name |
+-----+-----+
| 100        | Nokia       |
| 200        | Apple        |
| 300        | Samsung      |
+-----+-----+
```

Result table:

```
+-----+-----+-----+-----+
| product_id | first_year | quantity | price |
+-----+-----+-----+-----+
| 100        | 2008      | 10       | 5000  |
| 200        | 2011      | 15       | 9000  |
+-----+-----+-----+-----+
```

1075. Project Employees I

Table: Project

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| project_id  | int    |
| employee_id | int    |
+-----+-----+
(project_id, employee_id) is the primary key of this table.
employee_id is a foreign key to Employee table.
Table: Employee
```

```
+-----+-----+
| Column Name      | Type   |
+-----+-----+
| employee_id      | int    |
| name             | varchar|
| experience_years  | int    |
+-----+-----+
employee_id is the primary key of this table.
```

Write an SQL query that reports the average experience years of all the employees for each project, rounded to 2 digits.

The query result format is in the following example:

Project table:

```
+-----+-----+
| project_id | employee_id |
+-----+-----+
| 1          | 1           |
| 1          | 2           |
| 1          | 3           |
| 2          | 1           |
| 2          | 4           |
+-----+-----+
```

Employee table:

```
+-----+-----+-----+
| employee_id | name   | experience_years |
+-----+-----+-----+
| 1           | Khaled | 3                |
| 2           | Ali    | 2                |
```

3	John	1	
4	Doe	2	
+-----+-----+-----+			

Result table:

+-----+-----+		
project_id	average_years	
+-----+-----+		
1	2.00	
2	2.50	
+-----+-----+		

The average experience years for the first project is $(3 + 2 + 1) / 3 = 2.00$ and for the second project is $(3 + 2) / 2 = 2.50$

1076. Project Employees II

Table: Project

+-----+-----+		
Column Name	Type	
+-----+-----+		
project_id	int	
employee_id	int	
+-----+-----+		

(project_id, employee_id) is the primary key of this table.

employee_id is a foreign key to Employee table.

Table: Employee

+-----+-----+		
Column Name	Type	
+-----+-----+		
employee_id	int	
name	varchar	
experience_years	int	
+-----+-----+		

employee_id is the primary key of this table.

Write an SQL query that reports all the projects that have the most employees.

The query result format is in the following example:

Project table:

+-----+-----+		
---------------	--	--

project_id	employee_id
1	1
1	2
1	3
2	1
2	4

Employee table:

employee_id	name	experience_years
1	Khaled	3
2	Ali	2
3	John	1
4	Doe	2

Result table:

project_id
1

The first project has 3 employees while the second one has 2.

1077. Project Employees III

Table: Project

Column Name	Type
project_id	int
employee_id	int

(project_id, employee_id) is the primary key of this table.
employee_id is a foreign key to Employee table.

Table: Employee

Column Name	Type
employee_id	int

```

| name          | varchar |
| experience_years | int     |
+-----+-----+
employee_id is the primary key of this table.

```

Write an SQL query that reports the most experienced employees in each project. In case of a tie, report all employees with the maximum number of experience years.

The query result format is in the following example:

Project table:

```

+-----+-----+
| project_id | employee_id |
+-----+-----+
| 1          | 1           |
| 1          | 2           |
| 1          | 3           |
| 2          | 1           |
| 2          | 4           |
+-----+-----+

```

Employee table:

```

+-----+-----+-----+
| employee_id | name    | experience_years |
+-----+-----+-----+
| 1          | Khaled | 3                |
| 2          | Ali    | 2                |
| 3          | John   | 3                |
| 4          | Doe    | 2                |
+-----+-----+-----+

```

Result table:

```

+-----+-----+
| project_id | employee_id |
+-----+-----+
| 1          | 1           |
| 1          | 3           |
| 2          | 1           |
+-----+-----+

```

Both employees with id 1 and 3 have the most experience among the employees of the first project. For the second project, the employee with id 1 has the most experience.

1082. Sales Analysis I

Table: Product

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product_id is the primary key of this table.

Table: Sales

Column Name	Type
seller_id	int
product_id	int
buyer_id	int
sale_date	date
quantity	int
price	int

This table has no primary key, it can have repeated rows.

product_id is a foreign key to the Product table.

Write an SQL query that reports the best seller by total sales price, If there is a tie, report them all.

The query result format is in the following example:

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
-----------	------------	----------	-----------	----------	-------

1	1	1	2019-01-21	2	2000	
1	2	2	2019-02-17	1	800	
2	2	3	2019-06-02	1	800	
3	3	4	2019-05-13	2	2800	
+-----+-----+-----+-----+-----+-----+						

Result table:

+-----+	
seller_id	
+-----+	
1	
3	
+-----+	

Both sellers with id 1 and 3 sold products with the most total price of 2800.

1083. Sales Analysis II

Table: Product

+-----+-----+		
Column Name	Type	
+-----+-----+		
product_id	int	
product_name	varchar	
unit_price	int	
+-----+-----+		

product_id is the primary key of this table.

Table: Sales

+-----+-----+		
Column Name	Type	
+-----+-----+		
seller_id	int	
product_id	int	
buyer_id	int	
sale_date	date	
quantity	int	
price	int	
+-----+-----+		

This table has no primary key, it can have repeated rows.

product_id is a foreign key to Product table.

Write an SQL query that reports the buyers who have bought S8 but not iPhone. Note that S8 and iPhone are products present in the Product table.

The query result format is in the following example:

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	1	3	2019-06-02	1	800
3	3	3	2019-05-13	2	2800

Result table:

buyer_id
1

The buyer with id 1 bought an S8 but didn't buy an iPhone. The buyer with id 3 bought both.

1084. Sales Analysis III

Table: Product

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product_id is the primary key of this table.

Table: Sales

Column Name	Type
seller_id	int
product_id	int
buyer_id	int
sale_date	date
quantity	int
price	int

This table has no primary key, it can have repeated rows.
product_id is a foreign key to Product table.

Write an SQL query that reports the products that were only sold in spring 2019. That is, between 2019-01-01 and 2019-03-31 inclusive.

The query result format is in the following example:

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

Result table:

product_id	product_name
1	S8

The product with id 1 was only sold in spring 2019 while the other two were sold after.

1112. Highest Grade For Each Student

Table: Enrollments

Column Name	Type
student_id	int
course_id	int
grade	int

(student_id, course_id) is the primary key of this table.

Write a SQL query to find the highest grade with its corresponding course for each student. In case of a tie, you should find the course with the smallest course_id. The output must be sorted by increasing student_id.

The query result format is in the following example:

Enrollments table:

student_id	course_id	grade
2	2	95
2	3	95
1	1	90
1	2	99
3	1	80
3	2	75
3	3	82

Result table:

student_id	course_id	grade
1	2	99
2	2	95
3	3	82

1113. Reported Posts

Table: Actions

Column Name	Type
user_id	int
post_id	int
action_date	date
action	enum
extra	varchar

There is no primary key for this table, it may have duplicate rows.

The action column is an ENUM type of ('view', 'like', 'reaction', 'comment', 'report', 'share').

The extra column has optional information about the action such as a reason for report or a type of reaction.

Write an SQL query that reports the number of posts reported yesterday for each report reason. Assume today is 2019-07-05.

The query result format is in the following example:

Actions table:

user_id	post_id	action_date	action	extra
1	1	2019-07-01	view	null
1	1	2019-07-01	like	null
1	1	2019-07-01	share	null
2	4	2019-07-04	view	null
2	4	2019-07-04	report	spam
3	4	2019-07-04	view	null
3	4	2019-07-04	report	spam
4	3	2019-07-02	view	null
4	3	2019-07-02	report	spam
5	2	2019-07-04	view	null
5	2	2019-07-04	report	racism
5	5	2019-07-04	view	null
5	5	2019-07-04	report	racism

Result table:

|--|--|

report_reason	report_count
spam	1
racism	2

Note that we only care about report reasons with non zero number of reports.

1126. Active Businesses

Table: Events

Column Name	Type
business_id	int
event_type	varchar
occurences	int

(business_id, event_type) is the primary key of this table.

Each row in the table logs the info that an event of some type occurred at some business for a number of times.

Write an SQL query to find all active businesses.

An active business is a business that has more than one event type with occurrences greater than the average occurrences of that event type among all businesses.

The query result format is in the following example:

Events table:

business_id	event_type	occurences
1	reviews	7
3	reviews	3
1	ads	11
2	ads	7
3	ads	6
1	page views	3
2	page views	12

Result table:

```
+-----+
| business_id |
+-----+
| 1           |
+-----+
```

Average for 'reviews', 'ads' and 'page views' are $(7+3)/2=5$, $(11+7+6)/3=8$, $(3+12)/2=7.5$ respectively.

Business with id 1 has 7 'reviews' events (more than 5) and 11 'ads' events (more than 8) so it is an active business.

1141. User Activity for the Past 30 Days I

Table: Activity

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| user_id     | int    |
| session_id  | int    |
| activity_date | date   |
| activity_type | enum   |
+-----+-----+
```

There is no primary key for this table, it may have duplicate rows.

The activity_type column is an ENUM of type ('open_session', 'end_session', 'scroll_down', 'send_message').

The table shows the user activities for a social media website.

Note that each session belongs to exactly one user.

Write an SQL query to find the daily active user count for a period of 30 days ending 2019-07-27 inclusively. A user was active on some day if he/she made at least one activity on that day.

The query result format is in the following example:

Activity table:

```
+-----+-----+-----+-----+
| user_id | session_id | activity_date | activity_type |
+-----+-----+-----+-----+
| 1       | 1          | 2019-07-20    | open_session  |
| 1       | 1          | 2019-07-20    | scroll_down    |
| 1       | 1          | 2019-07-20    | end_session   |
| 2       | 4          | 2019-07-20    | open_session  |
| 2       | 4          | 2019-07-21    | send_message  |
```

2	4	2019-07-21	end_session	
3	2	2019-07-21	open_session	
3	2	2019-07-21	send_message	
3	2	2019-07-21	end_session	
4	3	2019-06-25	open_session	
4	3	2019-06-25	end_session	

+-----+-----+-----+-----+

Result table:

+-----+	+-----+
day	active_users
+-----+	+-----+
2019-07-20	2
2019-07-21	2
+-----+	+-----+

Note that we do not care about days with zero active users.

1142. User Activity for the Past 30 Days II

Table: Activity

+-----+	+-----+
Column Name	Type
+-----+	+-----+
user_id	int
session_id	int
activity_date	date
activity_type	enum
+-----+	+-----+

There is no primary key for this table, it may have duplicate rows.

The activity_type column is an ENUM of type ('open_session', 'end_session', 'scroll_down', 'send_message').

The table shows the user activities for a social media website.

Note that each session belongs to exactly one user.

Write an SQL query to find the average number of sessions per user for a period of 30 days ending 2019-07-27 inclusively, rounded to 2 decimal places. The sessions we want to count for a user are those with at least one activity in that time period.

The query result format is in the following example:

Activity table:

+-----+	+-----+	+-----+	+-----+
---------	---------	---------	---------

user_id	session_id	activity_date	activity_type
1	1	2019-07-20	open_session
1	1	2019-07-20	scroll_down
1	1	2019-07-20	end_session
2	4	2019-07-20	open_session
2	4	2019-07-21	send_message
2	4	2019-07-21	end_session
3	2	2019-07-21	open_session
3	2	2019-07-21	send_message
3	2	2019-07-21	end_session
3	5	2019-07-21	open_session
3	5	2019-07-21	scroll_down
3	5	2019-07-21	end_session
4	3	2019-06-25	open_session
4	3	2019-06-25	end_session

Result table:

average_sessions_per_user
1.33

User 1 and 2 each had 1 session in the past 30 days while user 3 had 2 sessions so the average is $(1 + 1 + 2) / 3 = 1.33$.

1148. Article Views I

Table: Views

Column Name	Type
article_id	int
author_id	int
viewer_id	int
view_date	date

There is no primary key for this table, it may have duplicate rows. Each row of this table indicates that some viewer viewed an article (written by some author) on some date. Note that equal author_id and viewer_id indicate the same person.

Write an SQL query to find all the authors that viewed at least one of their own articles, sorted in ascending order by their id.

The query result format is in the following example:

Views table:

article_id	author_id	viewer_id	view_date
1	3	5	2019-08-01
1	3	6	2019-08-02
2	7	7	2019-08-01
2	7	6	2019-08-02
4	7	1	2019-07-22
3	4	4	2019-07-21
3	4	4	2019-07-21

Result table:

id
4
7

1164. Product Price at a Given Date

Table: Products

Column Name	Type
product_id	int
new_price	int
change_date	date

(product_id, change_date) is the primary key of this table.
Each row of this table indicates that the price of some product was changed to a new price at some date.

Write an SQL query to find the prices of all products on 2019-08-16.
Assume the price of all products before any change is 10.

The query result format is in the following example:

Products table:

product_id	new_price	change_date
1	20	2019-08-14
2	50	2019-08-14
1	30	2019-08-15
1	35	2019-08-16
2	65	2019-08-17
3	20	2019-08-18

Result table:

product_id	price
2	50
1	35
3	10

1173. Immediate Food Delivery I

Table: Delivery

Column Name	Type
delivery_id	int
customer_id	int
order_date	date
customer_pref_delivery_date	date

delivery_id is the primary key of this table.

The table holds information about food delivery to customers that make orders at some date and specify a preferred delivery date (on the same order date or after it).

If the preferred delivery date of the customer is the same as the order date then the order is called immediate otherwise it's called scheduled.

Write an SQL query to find the percentage of immediate orders in the table, rounded to 2 decimal places.

The query result format is in the following example:

Delivery table:

delivery_id	customer_id	order_date	customer_pref_delivery_date
1	1	2019-08-01	2019-08-02
2	5	2019-08-02	2019-08-02
3	1	2019-08-11	2019-08-11
4	3	2019-08-24	2019-08-26
5	4	2019-08-21	2019-08-22
6	2	2019-08-11	2019-08-13

Result table:

immediate_percentage
33.33

The orders with delivery id 2 and 3 are immediate while the others are scheduled.

1174. Immediate Food Delivery II

Table: Delivery

Column Name	Type
delivery_id	int
customer_id	int
order_date	date
customer_pref_delivery_date	date

delivery_id is the primary key of this table.

The table holds information about food delivery to customers that make orders at some date and specify a preferred delivery date (on the same order date or after it).

If the preferred delivery date of the customer is the same as the order date then the order is called immediate otherwise it's called scheduled.

The first order of a customer is the order with the earliest order date that customer made. It is guaranteed that a customer has exactly one first order.

Write an SQL query to find the percentage of immediate orders in the first orders of all customers, rounded to 2 decimal places.

The query result format is in the following example:

Delivery table:

delivery_id	customer_id	order_date	customer_pref_delivery_date
1	1	2019-08-01	2019-08-02
2	2	2019-08-02	2019-08-02
3	1	2019-08-11	2019-08-12
4	3	2019-08-24	2019-08-24
5	3	2019-08-21	2019-08-22
6	2	2019-08-11	2019-08-13
7	4	2019-08-09	2019-08-09

Result table:

immediate_percentage
50.00

The customer id 1 has a first order with delivery id 1 and it is scheduled.

The customer id 2 has a first order with delivery id 2 and it is immediate.

The customer id 3 has a first order with delivery id 5 and it is scheduled.

The customer id 4 has a first order with delivery id 7 and it is immediate.

Hence, half the customers have immediate first orders.

1179. Reformat Department Table

Table: Department

```

+-----+-----+
| Column Name | Type |
+-----+-----+
| id          | int  |
| revenue     | int  |
| month       | varchar |
+-----+-----+

```

(id, month) is the primary key of this table.
The table has information about the revenue of each department per month.
The month has values in
["Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"].

Write an SQL query to reformat the table such that there is a department id column and a revenue column for each month.

The query result format is in the following example:

Department table:

```

+-----+-----+-----+
| id | revenue | month |
+-----+-----+-----+
| 1 | 8000 | Jan |
| 2 | 9000 | Jan |
| 3 | 10000 | Feb |
| 1 | 7000 | Feb |
| 1 | 6000 | Mar |
+-----+-----+-----+

```

Result table:

```

+-----+-----+-----+-----+-----+-----+
| id | Jan_Revenue | Feb_Revenue | Mar_Revenue | ... | Dec_Revenue |
+-----+-----+-----+-----+-----+-----+
| 1 | 8000 | 7000 | 6000 | ... | null |
| 2 | 9000 | null | null | ... | null |
| 3 | null | 10000 | null | ... | null |
+-----+-----+-----+-----+-----+-----+

```

Note that the result table has 13 columns (1 for the department id + 12 for the months).

1193. Monthly Transactions I

Table: Transactions

Column Name	Type
id	int
country	varchar
state	enum
amount	int
trans_date	date

id is the primary key of this table.

The table has information about incoming transactions.

The state column is an enum of type ["approved", "declined"].

Write an SQL query to find for each month and country, the number of transactions and their total amount, the number of approved transactions and their total amount.

The query result format is in the following example:

Transactions table:

id	country	state	amount	trans_date
121	US	approved	1000	2018-12-18
122	US	declined	2000	2018-12-19
123	US	approved	2000	2019-01-01
124	DE	approved	2000	2019-01-07

Result table:

month	country	trans_count	approved_count	trans_total_amount	approved_total_amount
2018-12	US	2	1	3000	1000
2019-01	US	1	1	2000	2000
2019-01	DE	1	1	2000	2000

1204. Last Person to Fit in the Elevator

Table: Queue

Column Name	Type
person_id	int
person_name	varchar
weight	int
turn	int

person_id is the primary key column for this table.

This table has the information about all people waiting for an elevator.

The person_id and turn columns will contain all numbers from 1 to n, where n is the number of rows in the table.

The maximum weight the elevator can hold is 1000.

Write an SQL query to find the person_name of the last person who will fit in the elevator without exceeding the weight limit. It is guaranteed that the person who is first in the queue can fit in the elevator.

The query result format is in the following example:

Queue table

person_id	person_name	weight	turn
5	George Washington	250	1
3	John Adams	350	2
6	Thomas Jefferson	400	3
2	Will Johnliams	200	4
4	Thomas Jefferson	175	5
1	James Elephant	500	6

Result table

person_name
Thomas Jefferson

Queue table is ordered by turn in the example for simplicity.

In the example George Washington(id 5), John Adams(id 3) and Thomas Jefferson(id 6) will enter the elevator as their weight sum is $250 + 350 + 400 = 1000$.
 Thomas Jefferson(id 6) is the last person to fit in the elevator because he has the last turn in these three people.

1211. Queries Quality and Percentage

Table: Queries

Column Name	Type
query_name	varchar
result	varchar
position	int
rating	int

There is no primary key for this table, it may have duplicate rows.
 This table contains information collected from some queries on a database.
 The position column has a value from 1 to 500.
 The rating column has a value from 1 to 5. Query with rating less than 3 is a poor query.

We define query quality as:

The average of the ratio between query rating and its position.

We also define poor query percentage as:

The percentage of all queries with rating less than 3.

Write an SQL query to find each query_name, the quality and poor_query_percentage.

Both quality and poor_query_percentage should be rounded to 2 decimal places.

The query result format is in the following example:

Queries table:

query_name	result	position	rating
------------	--------	----------	--------

Dog	Golden Retriever	1	5	
Dog	German Shepherd	2	5	
Dog	Mule	200	1	
Cat	Shirazi	5	2	
Cat	Siamese	3	3	
Cat	Sphynx	7	4	
+-----+-----+-----+-----+				

Result table:

query_name	quality	poor_query_percentage	
Dog	2.50	33.33	
Cat	0.66	33.33	
+-----+-----+-----+			

Dog queries quality is $((5 / 1) + (5 / 2) + (1 / 200)) / 3 = 2.50$

Dog queries poor_query_percentage is $(1 / 3) * 100 = 33.33$

Cat queries quality equals $((2 / 5) + (3 / 3) + (4 / 7)) / 3 = 0.66$

Cat queries poor_query_percentage is $(1 / 3) * 100 = 33.33$

1212. Team Scores in Football Tournament

Table: Teams

Column Name	Type	
team_id	int	
team_name	varchar	
+-----+-----+		

team_id is the primary key of this table.

Each row of this table represents a single football team.

Table: Matches

Column Name	Type	
match_id	int	
host_team	int	
guest_team	int	
host_goals	int	
guest_goals	int	
+-----+-----+		

match_id is the primary key of this table.
Each row is a record of a finished match between two different teams.
Teams host_team and guest_team are represented by their IDs in the teams table (team_id) and they scored host_goals and guest_goals goals respectively.

You would like to compute the scores of all teams after all matches.
Points are awarded as follows:
A team receives three points if they win a match (Score strictly more goals than the opponent team).
A team receives one point if they draw a match (Same number of goals as the opponent team).
A team receives no points if they lose a match (Score less goals than the opponent team).
Write an SQL query that selects the team_id, team_name and num_points of each team in the tournament after all described matches. Result table should be ordered by num_points (decreasing order). In case of a tie, order the records by team_id (increasing order).

The query result format is in the following example:

Teams table:

team_id	team_name
10	Leetcode FC
20	NewYork FC
30	Atlanta FC
40	Chicago FC
50	Toronto FC

Matches table:

match_id	host_team	guest_team	host_goals	guest_goals
1	10	20	3	0
2	30	10	2	2
3	10	50	5	1
4	20	30	1	0
5	50	30	1	0

Result table:

team_id	team_name	num_points
---------	-----------	------------

10	Leetcode FC	7
20	NewYork FC	3
50	Toronto FC	3
30	Atlanta FC	1
40	Chicago FC	0

1225. Report Contiguous Dates

Table: Failed

Column Name	Type
fail_date	date

Primary key for this table is fail_date.

Failed table contains the days of failed tasks.

Table: Succeeded

Column Name	Type
success_date	date

Primary key for this table is success_date.

Succeeded table contains the days of succeeded tasks.

A system is running one task every day. Every task is independent of the previous tasks. The tasks can fail or succeed.

Write an SQL query to generate a report of period_state for each continuous interval of days in the period from 2019-01-01 to 2019-12-31.

period_state is 'failed' if tasks in this interval failed or 'succeeded' if tasks in this interval succeeded. Interval of days are retrieved as start_date and end_date.

Order result by start_date.

The query result format is in the following example:

Failed table:

```

+-----+
| fail_date      |
+-----+
| 2018-12-28     |
| 2018-12-29     |
| 2019-01-04     |
| 2019-01-05     |
+-----+

```

Succeeded table:

```

+-----+
| success_date   |
+-----+
| 2018-12-30     |
| 2018-12-31     |
| 2019-01-01     |
| 2019-01-02     |
| 2019-01-03     |
| 2019-01-06     |
+-----+

```

Result table:

```

+-----+-----+-----+
| period_state | start_date | end_date   |
+-----+-----+-----+
| succeeded    | 2019-01-01 | 2019-01-03 |
| failed       | 2019-01-04 | 2019-01-05 |
| succeeded    | 2019-01-06 | 2019-01-06 |
+-----+-----+-----+

```

The report ignored the system state in 2018 as we care about the system in the period 2019-01-01 to 2019-12-31.

From 2019-01-01 to 2019-01-03 all tasks succeeded and the system state was "succeeded".

From 2019-01-04 to 2019-01-05 all tasks failed and the system state was "failed".

From 2019-01-06 to 2019-01-06 all tasks succeeded and the system state was "succeeded".

1241. Number of Comments per Post

Table: Submissions

```

+-----+-----+

```

Column Name	Type
sub_id	int
parent_id	int

There is no primary key for this table, it may have duplicate rows.
Each row can be a post or comment on the post.
parent_id is null for posts.
parent_id for comments is sub_id for another post in the table.

Write an SQL query to find number of comments per each post.

Result table should contain post_id and its corresponding number_of_comments, and must be sorted by post_id in ascending order.

Submissions may contain duplicate comments. You should count the number of unique comments per post.

Submissions may contain duplicate posts. You should treat them as one post.

The query result format is in the following example:

Submissions table:

sub_id	parent_id
1	Null
2	Null
1	Null
12	Null
3	1
5	2
3	1
4	1
9	1
10	2
6	7

Result table:

post_id	number_of_comments
1	3
2	2

```
| 12      | 0      |
+-----+-----+
```

The post with id 1 has three comments in the table with id 3, 4 and 9. The comment with id 3 is repeated in the table, we counted it only once.
The post with id 2 has two comments in the table with id 5 and 10.
The post with id 12 has no comments in the table.
The comment with id 6 is a comment on a deleted post with id 7 so we ignored it.

1251. Average Selling Price

Table: Prices

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| product_id  | int    |
| start_date  | date   |
| end_date    | date   |
| price       | int    |
+-----+-----+
```

(product_id, start_date, end_date) is the primary key for this table.
Each row of this table indicates the price of the product_id in the period from start_date to end_date.
For each product_id there will be no two overlapping periods. That means there will be no two intersecting periods for the same product_id.

Table: UnitsSold

```
+-----+-----+
| Column Name | Type   |
+-----+-----+
| product_id  | int    |
| purchase_date | date   |
| units       | int    |
+-----+-----+
```

There is no primary key for this table, it may contain duplicates.
Each row of this table indicates the date, units and product_id of each product sold.

Write an SQL query to find the average selling price for each product.

average_price should be rounded to 2 decimal places.

The query result format is in the following example:

Prices table:

product_id	start_date	end_date	price
1	2019-02-17	2019-02-28	5
1	2019-03-01	2019-03-22	20
2	2019-02-01	2019-02-20	15
2	2019-02-21	2019-03-31	30

UnitsSold table:

product_id	purchase_date	units
1	2019-02-25	100
1	2019-03-01	15
2	2019-02-10	200
2	2019-03-22	30

Result table:

product_id	average_price
1	6.96
2	16.96

Average selling price = Total Price of Product / Number of products sold.

Average selling price for product 1 = $((100 * 5) + (15 * 20)) / 115 = 6.96$

Average selling price for product 2 = $((200 * 15) + (30 * 30)) / 230 = 16.96$

1264. Page Recommendations

Table: Friendship

Column Name	Type
user1_id	int
user2_id	int

+-----+-----+

(user1_id, user2_id) is the primary key for this table.

Each row of this table indicates that there is a friendship relation between user1_id and user2_id.

Table: Likes

+-----+-----+

Column Name	Type
-------------	------

+-----+-----+

user_id	int
---------	-----

page_id	int
---------	-----

+-----+-----+

(user_id, page_id) is the primary key for this table.

Each row of this table indicates that user_id likes page_id.

Write an SQL query to recommend pages to the user with user_id = 1 using the pages that your friends liked. It should not recommend pages you already liked.

Return result table in any order without duplicates.

The query result format is in the following example:

Friendship table:

+-----+-----+

user1_id	user2_id
----------	----------

+-----+-----+

1	2
---	---

1	3
---	---

1	4
---	---

2	3
---	---

2	4
---	---

2	5
---	---

6	1
---	---

+-----+-----+

Likes table:

+-----+-----+

user_id	page_id
---------	---------

+-----+-----+

1	88
---	----

2	23
---	----

3	24
---	----

4	56
---	----

5	11	
6	33	
2	77	
3	77	
6	88	
+-----+-----+		

Result table:

+-----+		
recommended_page		
+-----+		
23		
24		
56		
33		
77		
+-----+		

User one is friend with users 2, 3, 4 and 6.

Suggested pages are 23 from user 2, 24 from user 3, 56 from user 3 and 33 from user 6.

Page 77 is suggested from both user 2 and user 3.

Page 88 is not suggested because user 1 already likes it.

1270. All People Report to the Given Manager

Table: Employees

+-----+-----+		
Column Name	Type	
+-----+-----+		
employee_id	int	
employee_name	varchar	
manager_id	int	
+-----+-----+		

employee_id is the primary key for this table.

Each row of this table indicates that the employee with ID employee_id and name employee_name reports his work to his/her direct manager with manager_id

The head of the company is the employee with employee_id = 1.

Write an SQL query to find employee_id of all employees that directly or indirectly report their work to the head of the company.

The indirect relation between managers will not exceed 3 managers as the company is small.

Return result table in any order without duplicates.

The query result format is in the following example:

Employees table:

employee_id	employee_name	manager_id
1	Boss	1
3	Alice	3
2	Bob	1
4	Daniel	2
7	Luis	4
8	Jhon	3
9	Angela	8
77	Robert	1

Result table:

employee_id
2
77
4
7

The head of the company is the employee with employee_id 1.

The employees with employee_id 2 and 77 report their work directly to the head of the company.

The employee with employee_id 4 report his work indirectly to the head of the company 4 --> 2 --> 1.

The employee with employee_id 7 report his work indirectly to the head of the company 7 --> 4 --> 2 --> 1.

The employees with employee_id 3, 8 and 9 don't report their work to head of company directly or indirectly.

1280. Students and Examinations

Table: Students

```

+-----+-----+
| Column Name | Type |
+-----+-----+
| student_id  | int  |
| student_name | varchar |
+-----+-----+

```

student_id is the primary key for this table.

Each row of this table contains the ID and the name of one student in the school.

Table: Subjects

```

+-----+-----+
| Column Name | Type |
+-----+-----+
| subject_name | varchar |
+-----+-----+

```

subject_name is the primary key for this table.

Each row of this table contains the name of one subject in the school.

Table: Examinations

```

+-----+-----+
| Column Name | Type |
+-----+-----+
| student_id  | int  |
| subject_name | varchar |
+-----+-----+

```

There is no primary key for this table. It may contain duplicates.

Each student from the Students table takes every course from Subjects table.

Each row of this table indicates that a student with ID student_id attended the exam of subject_name.

Write an SQL query to find the number of times each student attended each exam.

Order the result table by student_id and subject_name.

The query result format is in the following example:

Students table:

```

+-----+-----+
| student_id | student_name |

```

1	Alice
2	Bob
13	John
6	Alex

Subjects table:

subject_name
Math
Physics
Programming

Examinations table:

student_id	subject_name
1	Math
1	Physics
1	Programming
2	Programming
1	Physics
1	Math
13	Math
13	Programming
13	Physics
2	Math
1	Math

Result table:

student_id	student_name	subject_name	attended_exams
1	Alice	Math	3
1	Alice	Physics	2
1	Alice	Programming	1
2	Bob	Math	1
2	Bob	Physics	0
2	Bob	Programming	1
6	Alex	Math	0
6	Alex	Physics	0
6	Alex	Programming	0
13	John	Math	1
13	John	Physics	1
13	John	Programming	1

The result table should contain all students and all subjects.
 Alice attended Math exam 3 times, Physics exam 2 times and Programming exam 1 time.
 Bob attended Math exam 1 time, Programming exam 1 time and didn't attend the Physics exam.
 Alex didn't attend any exam.
 John attended Math exam 1 time, Physics exam 1 time and Programming exam 1 time.

1285. Find the Start and End Number of Continuous Ranges

Table: Logs

```
+-----+-----+
| Column Name | Type |
+-----+-----+
| log_id      | int  |
+-----+-----+
```

id is the primary key for this table.
 Each row of this table contains the ID in a log Table.

Since some IDs have been removed from Logs. Write an SQL query to find the start and end number of continuous ranges in table Logs.

Order the result table by start_id.

The query result format is in the following example:

Logs table:

```
+-----+
| log_id |
+-----+
| 1      |
| 2      |
| 3      |
| 7      |
| 8      |
| 10     |
+-----+
```

Result table:

```
+-----+-----+
| start_id | end_id |
+-----+-----+
| 1        | 3      |
```

7	8
10	10

The result table should contain all ranges in table Logs.
 From 1 to 3 is contained in the table.
 From 4 to 6 is missing in the table
 From 7 to 8 is contained in the table.
 Number 9 is missing in the table.
 Number 10 is contained in the table.

1294. Weather Type in Each Country

Table: Countries

Column Name	Type
country_id	int
country_name	varchar

country_id is the primary key for this table.
 Each row of this table contains the ID and the name of one country.

Table: Weather

Column Name	Type
country_id	int
weather_state	varchar
day	date

(country_id, day) is the primary key for this table.
 Each row of this table indicates the weather state in a country for one day.

Write an SQL query to find the type of weather in each country for November 2019.

The type of weather is Cold if the average weather_state is less than or equal 15, Hot if the average weather_state is greater than or equal 25 and Warm otherwise.

Return result table in any order.

The query result format is in the following example:

Countries table:

country_id	country_name
2	USA
3	Australia
7	Peru
5	China
8	Morocco
9	Spain

Weather table:

country_id	weather_state	day
2	15	2019-11-01
2	12	2019-10-28
2	12	2019-10-27
3	-2	2019-11-10
3	0	2019-11-11
3	3	2019-11-12
5	16	2019-11-07
5	18	2019-11-09
5	21	2019-11-23
7	25	2019-11-28
7	22	2019-12-01
7	20	2019-12-02
8	25	2019-11-05
8	27	2019-11-15
8	31	2019-11-25
9	7	2019-10-23
9	3	2019-12-23

Result table:

country_name	weather_type
USA	Cold
Austraila	Cold
Peru	Hot
China	Warm
Morocco	Hot

Average weather_state in USA in November is $(15) / 1 = 15$ so weather type is Cold.

Average weather_state in Austraila in November is $(-2 + 0 + 3) / 3 = 0.333$ so weather type is Cold.

Average weather_state in Peru in November is $(25) / 1 = 25$ so weather type is Hot.

Average weather_state in China in November is $(16 + 18 + 21) / 3 = 18.333$ so weather type is Warm.

Average weather_state in Morocco in November is $(25 + 27 + 31) / 3 = 27.667$ so weather type is Hot.

We know nothing about average weather_state in Spain in November so we don't include it in the result table.

1303. Find the Team Size

Table: Employee

Column Name	Type
employee_id	int
team_id	int

employee_id is the primary key for this table.

Each row of this table contains the ID of each employee and their respective team.

Write an SQL query to find the team size of each of the employees.

Return result table in any order.

The query result format is in the following example:

Employee Table:

employee_id	team_id
1	8
2	8
3	8
4	7
5	9
6	9

Result table:

employee_id	team_id
-------------	---------

employee_id	team_size
1	3
2	3
3	3
4	1
5	2
6	2

Employees with Id 1,2,3 are part of a team with team_id = 8.

Employees with Id 4 are part of a team with team_id = 7.

Employees with Id 5,6 are part of a team with team_id = 9.

1308. Running Total for Different Genders

Table: Scores

Column Name	Type
player_name	varchar
gender	varchar
day	date
score_points	int

(gender, day) is the primary key for this table.

A competition is held between the female team and males team.

Each row of this table indicates that a player_name and with gender has scored score_point in someday.

Gender is 'F' if the player is in the female team and 'M' if the player is in males team.

Write an SQL query to find the total score for each gender on each day.

Order the result table by gender and day

The query result format is in the following example:

Scores table:

player_name	gender	day	score_points
Aron	F	2020-01-01	17
Alice	F	2020-01-07	23

Bajrang	M	2020-01-07	7	
Khali	M	2019-12-25	11	
Slaman	M	2019-12-30	13	
Joe	M	2019-12-31	3	
Jose	M	2019-12-18	2	
Priya	F	2019-12-31	23	
Priyanka	F	2019-12-30	17	

+-----+-----+-----+-----+

Result table:

+-----+	+-----+	+-----+
gender	day	total
+-----+	+-----+	+-----+
F	2019-12-30	17
F	2019-12-31	40
F	2020-01-01	57
F	2020-01-07	80
M	2019-12-18	2
M	2019-12-25	13
M	2019-12-30	26
M	2019-12-31	29
M	2020-01-07	36
+-----+	+-----+	+-----+

For females team:

First day is 2019-12-30, Priyanka scored 17 points and the total score for the team is 17.

Second day is 2019-12-31, Priya scored 23 points and the total score for the team is 40.

Third day is 2020-01-01, Aron scored 17 points and the total score for the team is 57.

Fourth day is 2020-01-07, Alice scored 23 points and the total score for the team is 80.

For males team:

First day is 2019-12-18, Jose scored 2 points and the total score for the team is 2.

Second day is 2019-12-25, Khali scored 11 points and the total score for the team is 13.

Third day is 2019-12-30, Slaman scored 13 points and the total score for the team is 26.

Fourth day is 2019-12-31, Joe scored 3 points and the total score for the team is 29.

Fifth day is 2020-01-07, Bajrang scored 7 points and the total score for the team is 36.

1321. Restaurant Growth

Table: Customer

Column Name	Type
customer_id	int
name	varchar
visited_on	date
amount	int

(customer_id, visited_on) is the primary key for this table.
This table contains data about customer transactions in a restaurant.
visited_on is the date on which the customer with ID (customer_id) have visited the restaurant.
amount is the total paid by a customer.

You are the restaurant owner and you want to analyze a possible expansion (there will be at least one customer every day).

Write an SQL query to compute moving average of how much customer paid in a 7 days window (current day + 6 days before) .

The query result format is in the following example:

Return result table ordered by visited_on.

average_amount should be rounded to 2 decimal places, all dates are in the format ('YYYY-MM-DD').

Customer table:

customer_id	name	visited_on	amount
1	Jhon	2019-01-01	100
2	Daniel	2019-01-02	110
3	Jade	2019-01-03	120
4	Khaled	2019-01-04	130
5	Winston	2019-01-05	110
6	Elvis	2019-01-06	140
7	Anna	2019-01-07	150
8	Maria	2019-01-08	80
9	Jaze	2019-01-09	110
1	Jhon	2019-01-10	130
3	Jade	2019-01-10	150

+-----+-----+-----+-----+

Result table:

visited_on	amount	average_amount
2019-01-07	860	122.86
2019-01-08	840	120
2019-01-09	840	120
2019-01-10	1000	142.86

1st moving average from 2019-01-01 to 2019-01-07 has an average_amount of $(100 + 110 + 120 + 130 + 110 + 140 + 150)/7 = 122.86$

2nd moving average from 2019-01-02 to 2019-01-08 has an average_amount of $(110 + 120 + 130 + 110 + 140 + 150 + 80)/7 = 120$

3rd moving average from 2019-01-03 to 2019-01-09 has an average_amount of $(120 + 130 + 110 + 140 + 150 + 80 + 110)/7 = 120$

4th moving average from 2019-01-04 to 2019-01-10 has an average_amount of $(130 + 110 + 140 + 150 + 80 + 110 + 130 + 150)/7 = 142.86$

1322. Ads Performance

Table: Ads

Column Name	Type
ad_id	int
user_id	int
action	enum

(ad_id, user_id) is the primary key for this table.

Each row of this table contains the ID of an Ad, the ID of a user and the action taken by this user regarding this Ad.

The action column is an ENUM type of ('Clicked', 'Viewed', 'Ignored').

A company is running Ads and wants to calculate the performance of each Ad.

Performance of the Ad is measured using Click-Through Rate (CTR) where:

CTR = 0, if Ad total clicks + Ad total views = 0

CTR = Ad total clicks / (Ad total clicks + Ad total views) * 100, otherwise.

Write an SQL query to find the ctr of each Ad.

Round ctr to 2 decimal points. Order the result table by ctr in descending order and by ad_id in ascending order in case of a tie.

The query result format is in the following example:

Ads table:

ad_id	user_id	action
1	1	Clicked
2	2	Clicked
3	3	Viewed
5	5	Ignored
1	7	Ignored
2	7	Viewed
3	5	Clicked
1	4	Viewed
2	11	Viewed
1	2	Clicked

Result table:

ad_id	ctr
1	66.67
3	50.00
2	33.33
5	0.00

for ad_id = 1, ctr = $(2/(2+1)) * 100 = 66.67$

for ad_id = 2, ctr = $(1/(1+2)) * 100 = 33.33$

for ad_id = 3, ctr = $(1/(1+1)) * 100 = 50.00$

for ad_id = 5, ctr = 0.00, Note that ad_id = 5 has no clicks or views.

Note that we don't care about Ignored Ads.

Result table is ordered by the ctr. in case of a tie we order them by ad_id

1327. List the Products Ordered in a Period

Table: Products

```

+-----+-----+
| Column Name      | Type      |
+-----+-----+
| product_id       | int       |
| product_name     | varchar   |
| product_category | varchar   |
+-----+-----+

```

product_id is the primary key for this table.

This table contains data about the company's products.

Table: Orders

```

+-----+-----+
| Column Name      | Type      |
+-----+-----+
| product_id       | int       |
| order_date       | date      |
| unit             | int       |
+-----+-----+

```

There is no primary key for this table. It may have duplicate rows.

product_id is a foreign key to Products table.

unit is the number of products ordered in order_date.

Write an SQL query to get the names of products with greater than or equal to 100 units ordered in February 2020 and their amount.

Return result table in any order.

The query result format is in the following example:

Products table:

```

+-----+-----+-----+
| product_id | product_name          | product_category |
+-----+-----+-----+
| 1          | Leetcode Solutions    | Book             |
| 2          | Jewels of Stringology | Book             |
| 3          | HP                    | Laptop           |
| 4          | Lenovo                | Laptop           |
| 5          | Leetcode Kit          | T-shirt          |
+-----+-----+-----+

```

Orders table:

```

+-----+-----+-----+
| product_id | order_date | unit |
+-----+-----+-----+

```

1	2020-02-05	60	
1	2020-02-10	70	
2	2020-01-18	30	
2	2020-02-11	80	
3	2020-02-17	2	
3	2020-02-24	3	
4	2020-03-01	20	
4	2020-03-04	30	
4	2020-03-04	60	
5	2020-02-25	50	
5	2020-02-27	50	
5	2020-03-01	50	
+-----+-----+-----+			

Result table:

+-----+-----+	
product_name	unit
+-----+-----+	
Leetcode Solutions	130
Leetcode Kit	100
+-----+-----+	

Products with product_id = 1 is ordered in February a total of (60 + 70) = 130.

Products with product_id = 2 is ordered in February a total of 80.

Products with product_id = 3 is ordered in February a total of (2 + 3) = 5.

Products with product_id = 4 was not ordered in February 2020.

Products with product_id = 5 is ordered in February a total of (50 + 50) = 100.

1336. Number of Transactions per Visit

Table: Visits

+-----+-----+	
Column Name	Type
+-----+-----+	
user_id	int
visit_date	date
+-----+-----+	

(user_id, visit_date) is the primary key for this table.

Each row of this table indicates that user_id has visited the bank in visit_date.

Table: Transactions

Column Name	Type
user_id	int
transaction_date	date
amount	int

There is no primary key for this table, it may contain duplicates.
Each row of this table indicates that user_id has done a transaction of amount in transaction_date.
It is guaranteed that the user has visited the bank in the transaction_date.(i.e The Visits table contains (user_id, transaction_date) in one row)

A bank wants to draw a chart of the number of transactions bank visitors did in one visit to the bank and the corresponding number of visitors who have done this number of transaction in one visit.

Write an SQL query to find how many users visited the bank and didn't do any transactions, how many visited the bank and did one transaction and so on.

The result table will contain two columns:

transactions_count which is the number of transactions done in one visit.
visits_count which is the corresponding number of users who did transactions_count in one visit to the bank.
transactions_count should take all values from 0 to max(transactions_count) done by one or more users.

Order the result table by transactions_count.

The query result format is in the following example:

Visits table:

user_id	visit_date
1	2020-01-01
2	2020-01-02
12	2020-01-01
19	2020-01-03
1	2020-01-02

2	2020-01-03	
1	2020-01-04	
7	2020-01-11	
9	2020-01-25	
8	2020-01-28	

+-----+-----+

Transactions table:

user_id	transaction_date	amount	
1	2020-01-02	120	
2	2020-01-03	22	
7	2020-01-11	232	
1	2020-01-04	7	
9	2020-01-25	33	
9	2020-01-25	66	
8	2020-01-28	1	
9	2020-01-25	99	

+-----+-----+

Result table:

transactions_count	visits_count	
0	4	
1	5	
2	0	
3	1	

+-----+-----+

* For transactions_count = 0, The visits (1, "2020-01-01"), (2, "2020-01-02"), (12, "2020-01-01") and (19, "2020-01-03") did no transactions so visits_count = 4.

* For transactions_count = 1, The visits (2, "2020-01-03"), (7, "2020-01-11"), (8, "2020-01-28"), (1, "2020-01-02") and (1, "2020-01-04") did one transaction so visits_count = 5.

* For transactions_count = 2, No customers visited the bank and did two transactions so visits_count = 0.

* For transactions_count = 3, The visit (9, "2020-01-25") did three transactions so visits_count = 1.

* For transactions_count >= 4, No customers visited the bank and did more than three transactions so we will stop at transactions_count = 3